

# Optimization: Stochastic Gradient Descent

Convolutional Neural Networks for Visual Recognition

<http://cs231n.github.io/>

---

Simon Bihel, [simon.bihel@ens-rennes.fr](mailto:simon.bihel@ens-rennes.fr)

Friday 7<sup>th</sup> July, 2017

COINSE Lab, KAIST

# Table of contents

1. Introduction
2. Optimization
3. Computing the gradient

# Introduction

---

## Last week with Jinhan

A **score function** mapping the raw image pixels to class scores, e.g.

$$f(x_i, W) = Wx_i$$

With  $W$  being a set of parameters.

A **loss function** measuring the quality (how the scores agree with training data) of a set of parameters, e.g. SVM

$$L = \frac{1}{N} \sum_i \sum_{j \neq y_i} [\max(0, f(x_i; W)_j - f(x_i; W)_{y_i} + 1)] + \alpha R(W)$$

We want to minimize the loss function by tweaking  $W$ .

# Visualizing the loss function

Difficult to visualize due to high-dimensional spaces. Gain some intuition by choosing a random  $W$  and moving in a random direction ( $L(W + aW_1)$ )




1 dim, 1 example

# Visualizing the loss function

Difficult to visualize due to high-dimensional spaces. Gain some intuition by choosing a random  $W$  and moving in a random direction ( $L(W + aW_1)$ ) or 2 directions ( $L(W + aW_1 + bW_2)$ ).




1 dim, 1 example




2 dim, 1 example

# Visualizing the loss function


Difficult to visualize due to high-dimensional spaces. Gain some intuition by choosing a random  $W$  and moving in a random direction ( $L(W + aW_1)$ ) or 2 directions ( $L(W + aW_1 + bW_2)$ ).



1 dim, 1 example



2 dim, 1 example




2 dim, 100 examples

# Piecewise-linear structure of the loss function

For a single example we have

$$L_i = \sum_{j \neq y_i} \left[ \max(0, w_j^T x_i - w_{y_i}^T x_i + 1) \right]$$



In our examples the loss function is convex but it is not in practice with neural networks, thus we need more general minimization techniques.



# Optimization

---

## Strategy #1: Random Search

Obviously not the most accurate method on its own. But it is a good starting point to iteratively refine it.

## Strategy #2: Random Local Search

Start from a random  $W$  and add a random perturbation  $W + \delta W$  and perform an update if the loss is lower.

Still wasteful and computationally expensive.

## Strategy #3: Following the gradient

No need to randomly search for a good direction: we can compute the *best* direction with the steepest descend. This direction is called the **gradient** of the loss function

The gradient is just a vector of slopes (i.e. **derivatives**)

For a 1-D function it is:

$$\frac{df(x)}{dx} = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$

For vectors it is simply the vector of derivatives for each dimension (called **partial derivatives**)

# Computing the gradient

---

# Numerically with finite differences

Slow, approximate but easy way.

For each dimension, make a small change and compute the partial derivative along that dimension.

Choosing the step size (aka *learning rate*) is one of the most important hyperparameter in training a NN. Large changes can miss minimums but are faster.

# Analytically with calculus

Derive a direct formula for the gradient. No approximations and very fast to compute but error-prone in the implementation.

Lets use the example of the SVM loss function for a single datapoint:

$$L_i = \sum_{j \neq y_i} \left[ \max(0, w_j^T x_i - w_{y_i}^T x_i + \Delta) \right]$$

$$\nabla_{w_{y_i}} L_i = - \left( \sum_{j \neq y_i} \mathbb{1}(w_j^T x_i - w_{y_i}^T x_i + \Delta > 0) \right) x_i$$

You can check the correctness by comparing with the numerical gradient (**gradient check**).

# Gradient descent

It is the procedure of repeatedly evaluating the gradient and then performing a parameter update. This is the core of all NN libraries.

For large training data, a common approach is to compute the gradient over **batches** of data. This works well because training data are correlated.

**Stochastic Gradient Descent (SGD)** is the process of computing the gradient multiple times for only one example.